

CPU Scheduling

Xin Liu

xl24j@fsu.edu

COP 4610 Operating Systems

Homework 2 Posted on Canvas


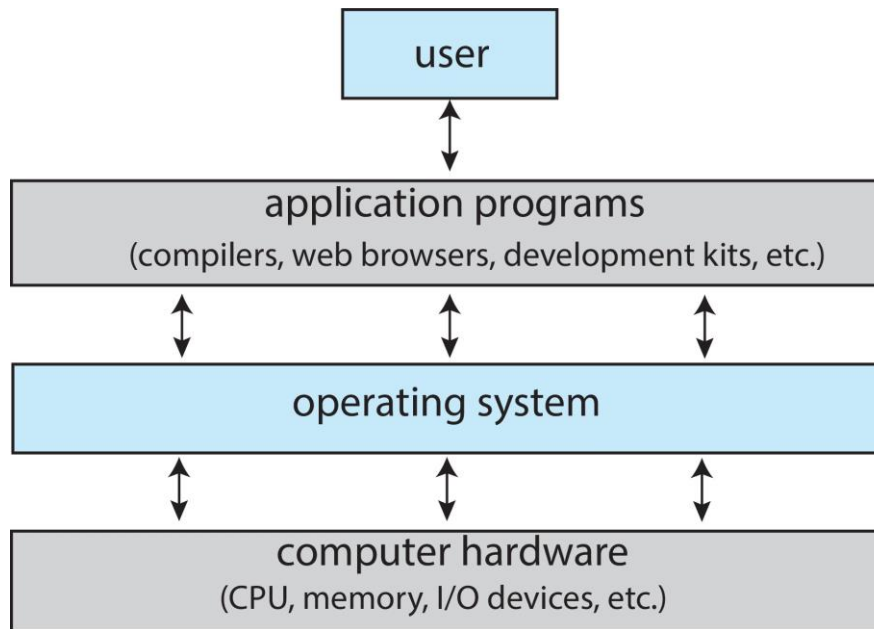
- **DDL: Mon, Sep 16th, 11:59 PM**

Outline

- How is the first process created?
- How is the second process created?
- How to manage processes?

Recap

- CPU is a state machine.
- A program (whether an OS or application) running on a CPU is inevitably a state machine.





If we are given the initial state, we can deduce what the next state will be.
The key is knowing:

Where is the initial state?

The Beginning of Everything

- An application runs on the OS
 - > OS creates the application's initial state



Who creates the initial state of OS?

Who creates the initial state of OS?



- Well...it's a long story...
 - It starts with a simple computing machine

Long, Long,
Long Ago...
(During the 1940s)

Long, Long, Long Ago...(during the 1940s)

- John von Neumann invented *von Neumann computer architecture*


- A CPU
- A memory unit
- I/O devices (e.g., disks and tapes)



von Neumann Architecture

In von Neumann Architecture,

- Programs are stored on storage devices
- Programs are copied into memory for execution
- CPU reads each instruction in the program and executes accordingly




A Simple CPU Model

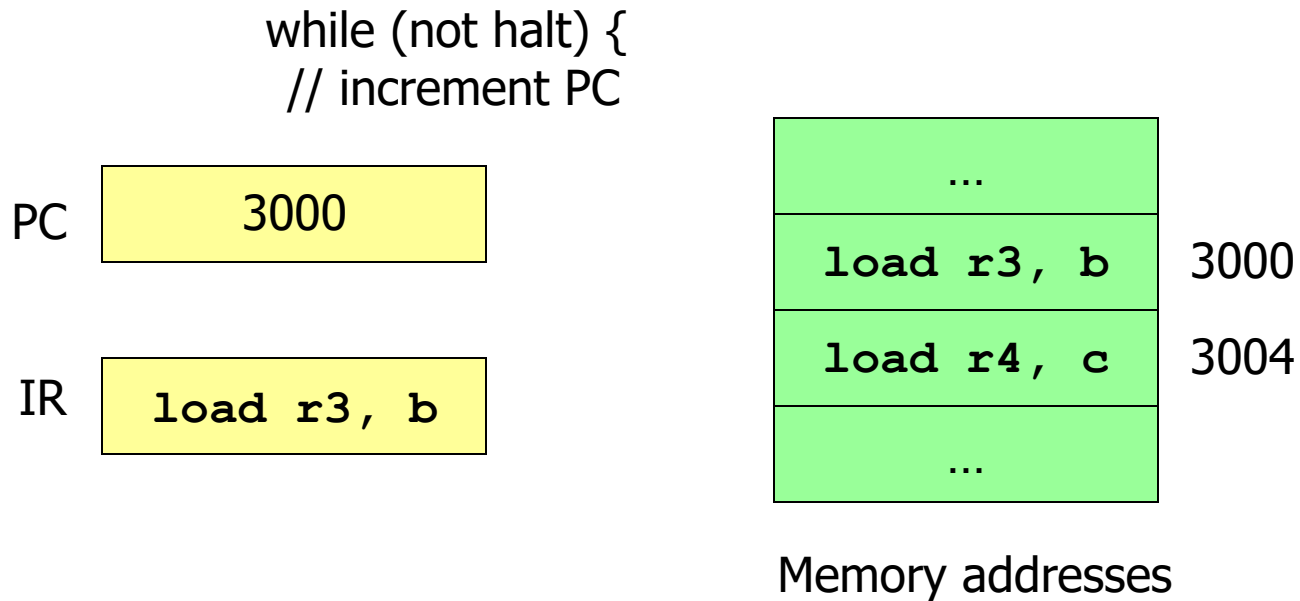
- *Fetch-execute algorithm*
- The CPU fetches the instruction:
 - The program counter (PC) is loaded with the address of the instruction
 - The instruction register (IR) is loaded with the instruction from the address
- The CPU decodes the instruction.
- The CPU executes the instruction.

The CPU fetches the instruction

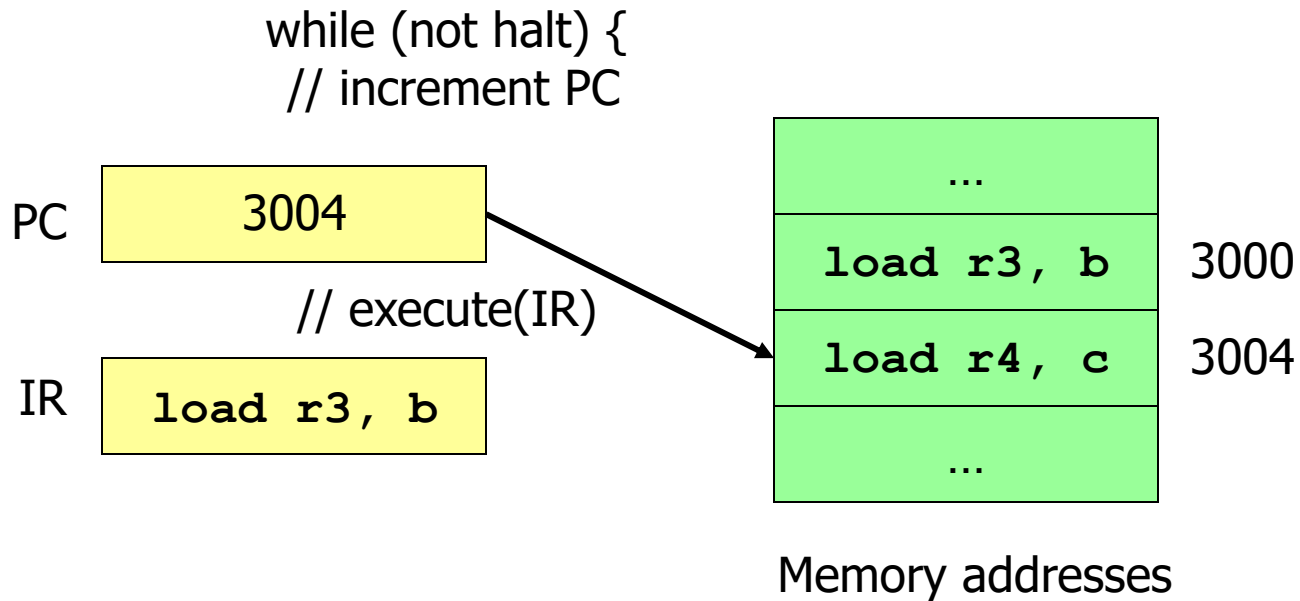
PC = <address of the first instruction>



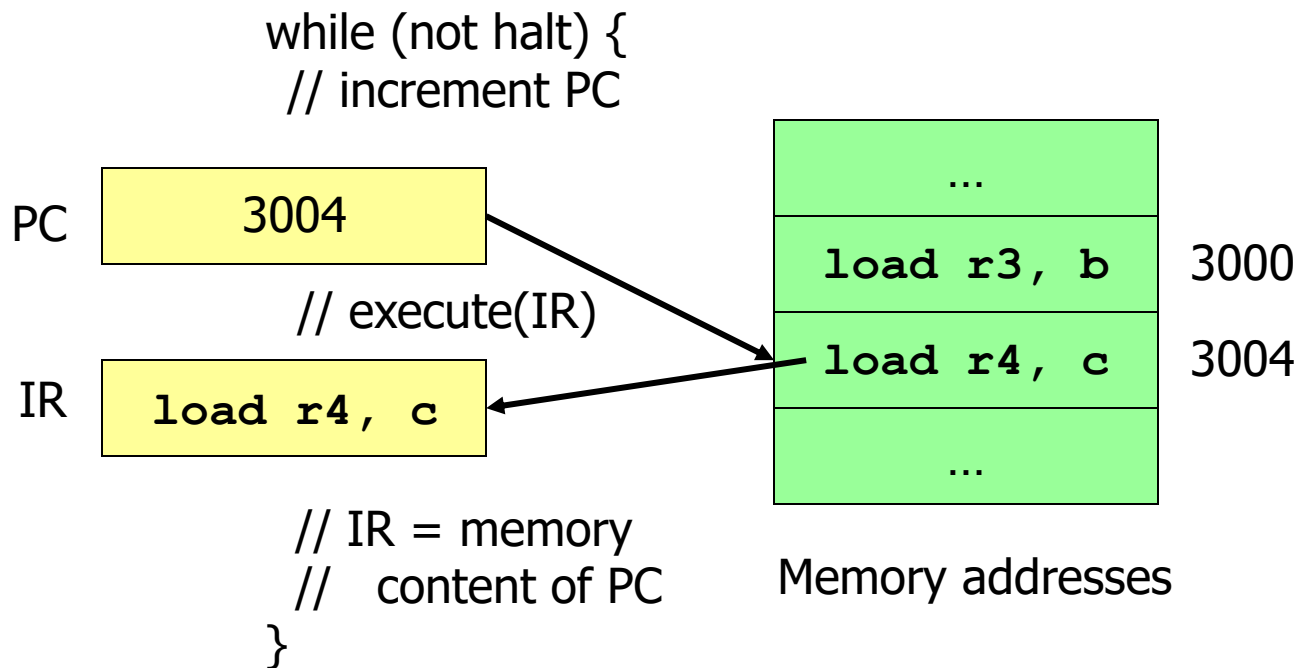
The CPU fetches the instruction



Fetch-Execute Algorithm



The CPU fetches the instruction



OS Booting Sequence

- The address of the first instruction is fixed
- It is stored in read-only-memory (ROM)

Booting Procedure for i386 Machines

- On i386 machines, ROM stores a *Basic Input/Output System (BIOS)*
 - BIOS contains information on how to access storage devices
- Being replaced with *United Extended Firmware Interface (UEFI)*
 - To access storage > 2TB

BIOS Code

- Performs Power-On Self Test (POST)
 - Checks memory and devices for their presence and correct operations
 - For ancient computers, you will hear memory counting, which consists of noises from the hard drive and CDROM, followed by a final beep

After the POST

- The *master boot record (MBR)* is loaded from the *boot device* (e.g., a hard drive configured in BIOS)
- The MBR is stored at the first logical sector of the boot device that
 - Fits into a single 512-byte disk sector (*boot sector*)
 - Describes the physical layout of the disk (e.g., number of tracks)
- MBR is being replaced by GUID Partition Table (*GPT*) for 64-bit addressing

After Getting the Info on the Boot Device

- BIOS loads a more sophisticated loader from other sectors on disk
 - Under old Linux, this sophisticated loader is called **LILLO** (*Linux Loader*)
 - It has nothing to do with Lilo and Stitch
 - Linux uses **GRUB** (*GRand Unified Bootloader*) nowadays
- The more sophisticated loader loads the operating system



More on OS Loaders

- LILO

- Partly stored in MBR with the disk partition table
 - A user can specify which disk partition and OS image to boot
 - Windows loader assumes only one bootable disk partition
- After loading the kernel image, LILO sets the kernel mode and jumps to the entry point of an operating system

Booting Sequence After A CPU Reset

- A CPU jumps to a fixed address in ROM,
- Loads the BIOS (UEFI),
- Performs POST,
- Loads MBR (GPT) from the boot device,
- Loads an OS loader (LILO, GRUB),
- Loads the kernel image,
- Sets the kernel mode, and
- Jumps to the OS entry point -- *init*

Read Latest Linux Kernel Code:

<https://elixir.bootlin.com/linux/v6.10.9/source/init/main.c#L1523>


Important

Can We Really See Every Instruction Executed After a CPU Reset?

- **"Talk is cheap. Show me the code."**— Linus Torvalds
- **Computer System Axiom:**
 - *If you can imagine it, someone has already done it.*
- **Simulation Option: QEMU**
 - Created by legendary hacker and genius programmer **Fabrice Bellard**
 - **QEMU**: A fast and portable dynamic translator (*USENIX ATC'05*)
 - Powers Android Virtual Devices, VirtualBox, and more
(*All built on QEMU*)
- **Real Machine Option: JTAG (Joint Test Action Group) Debugger**
 - A series of **physical debugging registers**
 - Allows integration with **gdb** (!!!)

A Side Story: Firmware Virus (1998)

- **Firmware is usually read-only** (though...)
 - The Intel 430TX (Pentium) chipset allows writing to Flash ROM.
 - By writing a specific sequence to Flash BIOS, the Flash ROM becomes writable.
 - This provides a channel for firmware updates.
 - Getting this sequence isn't too difficult.
 - It seems the documentation even provides it. 🙄 Boom...
- Chen Ing-Hau, the author of CIH, was arrested but not convicted.



Linux Initialization

- Set up a number of things:
 - **Trap table:** Handles system calls and exceptions.
 - **Interrupt handlers:** Manage external hardware interrupts.
 - **Scheduler:** Decides which processes run and when.
 - **Clock:** Manages system time and scheduling.
 - **Kernel modules:** Loads additional functionality for the kernel.
 - ...
 - **Process manager:** Controls process creation, termination, and state changes.

How is the first process created?

Process 1:

- Is instantiated from the *init* (now *systemd* for parallelism) program
- Is the ancestor of all processes
- Controls transitions between *runlevels*
- Executes startup and shutdown scripts for each runlevel

Runlevels

- Level 0: shutdown
- Level 1: single-user
- Level 2: multi-user (without network file system)
- Level 3: full multi-user
- Level 5: X11
- Level 6: reboot

Process Creation

- Via the *fork* system call family

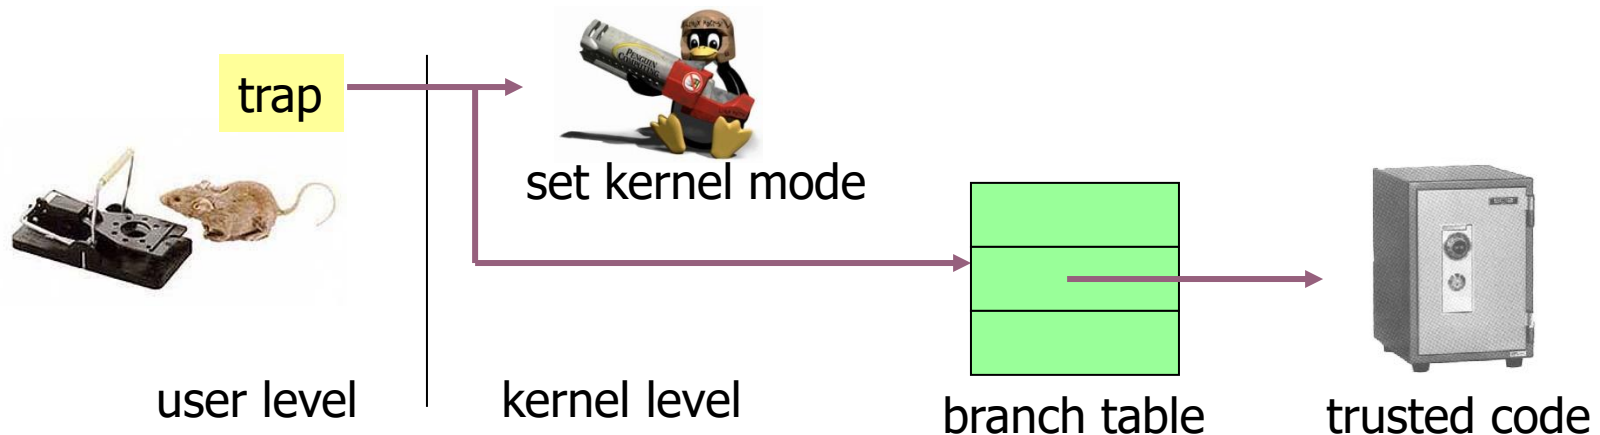
Before we discuss process creation, a few words on system calls...

System Calls

- *System calls* allow processes running at the *user mode* to access kernel functions that run under the *kernel mode*
- Prevent processes from doing bad things, such as
 - Halting the entire operating system
 - Modifying the MBR

UNIX System Calls

- Implemented through the *trap* instruction
 - The program in user mode makes a system call request (like read).
 - trap instruction is triggered, switching the CPU to kernel mode.
 - The kernel takes control and looks up the system call number.
 - It uses a branch table (system call table) to find the correct service.
 - The trusted code is executed to perform the requested operation.
 - The system then returns to user mode, and the program continues running.



Important

This Is How Simple Operating Systems Are!

CPU Reset → Firmware (BIOS/UEFI) → Boot Loader (MBR, LILO/GRUB)

↓
Kernel_start()

Use **ls /sbin/init -l** to check if systemd is being used

Read Latest Linux Kernel Code:

<https://elixir.bootlin.com/linux/v6.10.9/source/init/main.c#L1523>

↓
Process 1

↓
Application Program (state machine)
+

system call

↓
Process
Management

↓
Memory
Management

↓
File
Management

- Process Management
 - fork, exec, and exit
- Memory Management
 - mmap –virtual address space
- File management
 - open, close, read, write
 - mkdir, link, unlink

You can use system call to create the world!

fork()

- Creates a complete copy of the state machine
 - Copies memory, register states, and other process information
 - `int fork();` - system call used to create a new process
- How `fork()` Works:
 - Immediately copies the state machine, including the entire memory space
 - The newly created process (child) returns 0
 - The process that calls `fork()` (parent) returns the child process ID (PID)

A fork Example, Nag.c

```
#include <stdio.h>
#include <unistd.h>
#include <sys/types.h>

int main() {
    pid_t pid;
    if ((pid = fork()) == 0) {
        while (1) {
            printf("child's return value %d:  I want to play...\n", pid);
        }
    } else {
        while (1) {
            printf("parent's return value %d:  After the project...\n", pid);
        }
    }
    return 0;
}
```

A fork Example, Nag.c

```
#include <stdio.h>
#include <unistd.h>
#include <sys/types.h>

int main() {
    pid_t pid;
    if ((pid = 3128) == 0) {
        while (1) {
            printf("child's return value %d: I want to play...\n", pid);
        }
    } else {
        while (1) {
            printf("parent's\n");
        }
    }
    return 0;
}
```

Parent process

```
#include <stdio.h>
#include <unistd.h>
#include <sys/types.h>

int main() {
    pid_t pid;
    if ((pid = 0) == 0) {
        while (1) {
            printf("child's return value %d: I want to play...\n", pid);
        }
    } else {
        while (1) {
            printf("parent's return value %d: After the project...\n", pid);
        }
    }
    return 0;
}
```

Child process

Nag . c Outputs

>a.out

child's return value 0: I want to play...

child's return value 0: I want to play...

child's return value 0: I want to play...

...// context switch

parent's return value 3218: After the project...

parent's return value 3218: After the project...

parent's return value 3218: After the project...

...// context switch

child's return value 0: I want to play...

child's return value 0: I want to play...

child's return value 0: I want to play...

^C

>

Why clone a process?

Why clone a process?


- Simplifies parameter passing
 - Environmental variables, permissions, etc.
- Performance optimization
 - Copy on write

Also,...don't try it (or try it in docker)

Fork Bomb:

```
while (1) {  
    fork();  
}
```

Or



In bash, colons are allowed as identifiers...
the symbol for fork is a colon.

```
fork() {fork | fork &} ; fork
```

The **exec** System Call Family

- A **fork** by itself is not interesting
- To make a process run a program that is different from the parent process, you need **exec** system call
- **exec** starts a program by overwriting the current process

E.g.,

- `int execve(const char *filename, char * const argv[], char * const envp[]);`
 - Executes a program named filename
 - Allows setting parameters argv (arguments) and environment variables envp
 - Directly corresponds to the arguments passed to `main()`!

The **exit** System Call Family

- Destroy the current state machine and allow for a return value.
- The termination of the child process will notify the parent process (explained in later lessons)

Thread Creation

- Use **pthread_create()** instead of **fork()**
 - A newly created thread will share the address space of the current process and all resources (e.g., open files)
- + Efficient sharing of states
- Potential corruptions by a misbehaving thread

Address Space

- Contains all states necessary to run a program
 - Code, data, stack
 - Program counter
 - Register values
 - Resources required by the program
 - Status of the running program
- A mechanism to protect one app from crashing another app
- Inspect the address space of a process
 - pmap - report memory of a process

Understanding Address Space through Game Cheats

- Pioneers of esports: Real-Time Strategy (RTS) Games
- Command and Conquer (Westwood), Starcraft (Blizzard), ...
- What if we wanted to 'tamper with' the execution of the game...?
 - `ps aux | grep starcraft`
 - `pmap <pid>`
 - change data:
 - `gdb -p <process_id>`
 - `x /10x 0x<address>`
 - `set *(int *)0x<address> = <new_value>`



Understanding Address Space through Game Cheats

- Game is also state machine.

```
ps aux | grep pugb
```

```
pmap <pid>
```

```
change data:
```

```
gdb -p <process_id>
```

```
x /10x 0x<address>
```

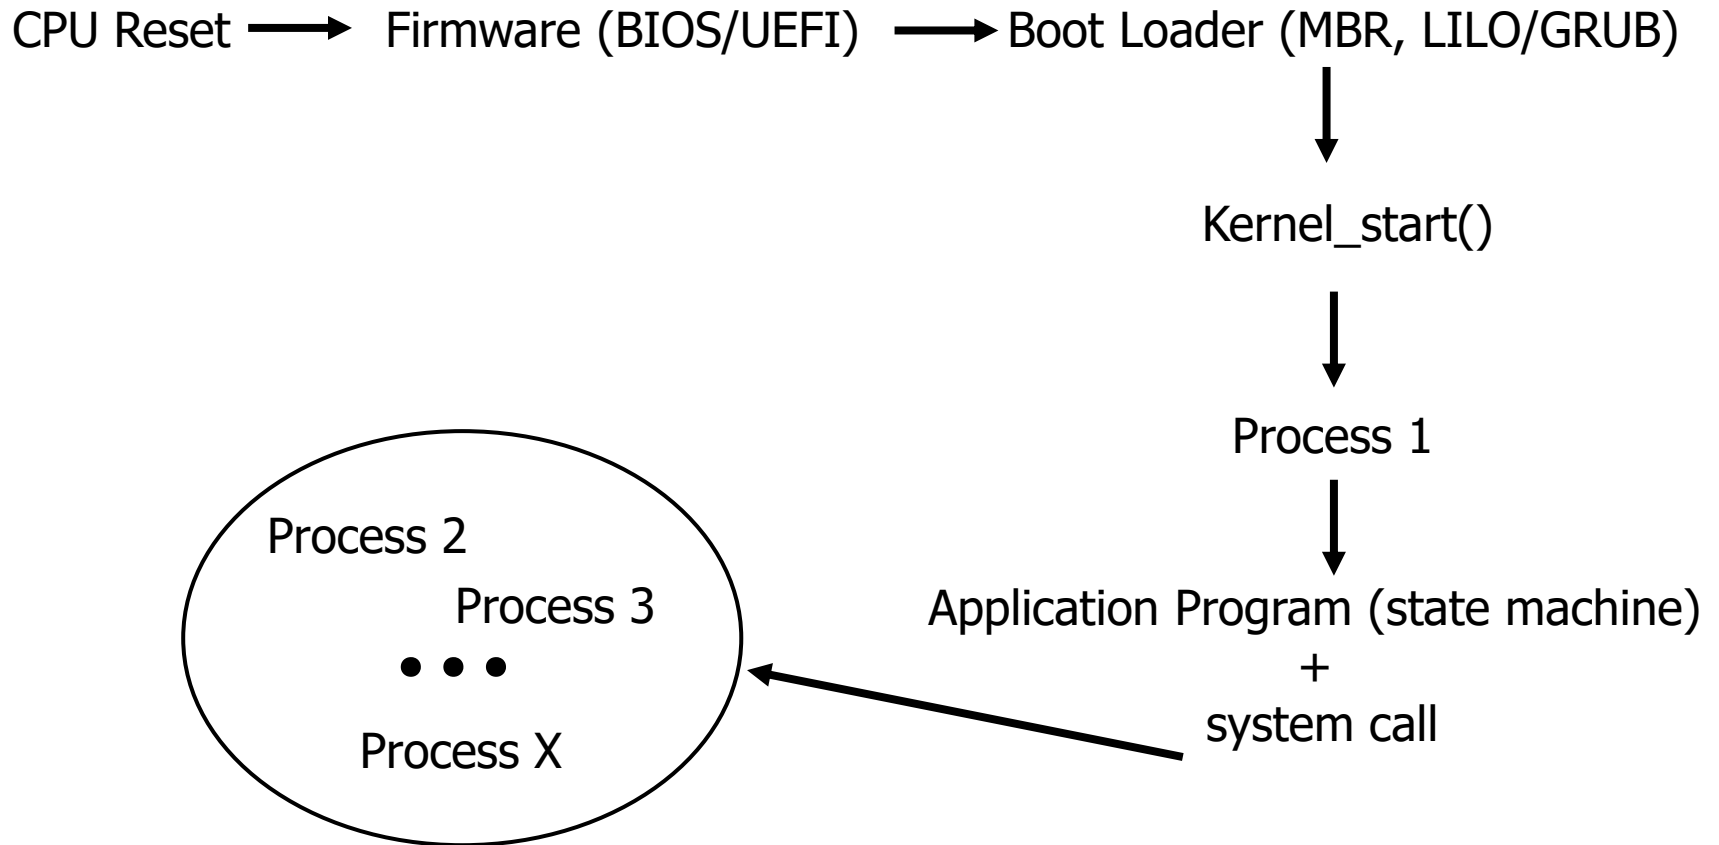
```
set *(int *)0x<address> = <new_value>
```



PUBG: BATTLEGROUNDS

CPU Scheduler

- A **CPU scheduler** is responsible for
 - Removal of running process from the CPU
 - Selection of the next running process
 - Based on a particular strategy



Goals for a Scheduler


- Maximize
 - *CPU utilization*: keep the CPU as busy as possible
 - *Throughput*: the number of processes completed per unit time

Goals for a Scheduler

- Minimize
 - *Response time*: the time of submission to the time the first response is produced
 - *Wait time*: total time spent waiting in the ready queue
 - *Turnaround time*: the time of submission to the time of completion


Goals for a Scheduler

- Suppose we have processes A, B, and C, submitted at time 0
- We want to know the response time, wait time, and turnaround time of process A




Goals for a Scheduler

- Suppose we have processes A, B, and C, submitted at time 0
- We want to know the response time, wait time, and turnaround time of process B




Goals for a Scheduler

- Suppose we have processes A, B, and C, submitted at time 0
- We want to know the response time, wait time, and turnaround time of process C



Goals for a Scheduler

- Suppose we have processes A and B submitted at time 0; process C, time 1
- We want to know the response time, wait time, and turnaround time of process C



Goals for a Scheduler

- Achieve *fairness*
 - What is fair?
 - Guaranteed to have at least $1/n$ share
 - How do two people divide a cake in a fair way?
- There are tensions among these goals

Assumptions

- Each user runs one process
- Each process is single threaded
- Processes are independent
- They are not realistic assumptions; they serve to simplify analyses

Scheduling Policies

- FIFO (first in, first out)
- Round robin
- SJF (shortest job first)
- Multilevel feedback queues
- Lottery scheduling

FIFO

- **FIFO**: assigns the CPU based on the order of requests
 - **Nonpreemptive**: A process keeps running on a CPU until it is blocked or terminated
 - Also known as FCFS (first come, first serve)
 - + Simple
 - Short jobs can get stuck behind long jobs


Round Robin

- **Round Robin (RR)** periodically releases the CPU from long-running jobs
 - Based on timer interrupts so short jobs can get a fair share of CPU time
 - **Preemptive**: a process can be forced to leave its running state and replaced by another running process
 - **Time slice**: interval between timer interrupts


More on Round Robin

- If time slice is too long
 - Scheduling degrades to FIFO
- If time slice is too short
 - Throughput suffers
 - Context switching cost dominates


Round Robin (Time slice = 100)




Round Robin (Time slice = 100)




Round Robin (Time slice = 100)




Round Robin (Time slice = 100)




Round Robin (Time slice = 100)




Round Robin (Time slice = 100)




Round Robin (Time slice = 100)




Round Robin (Time slice = 100)




Round Robin (Time slice = 100)




Round Robin (Time slice = 100)




Round Robin (Time slice = 100)




Round Robin (Time slice = 100)




Round Robin (Time slice = 100)




Round Robin (Time slice = 100)




Round Robin (Time slice = 100)




Round Robin (Time slice = 100)




Round Robin (Time slice = 100)



Round Robin (Time slice = 100)



Round Robin (Time slice = 100)




Response time for process 1: 0

Response time for process 2: $300 - 150 = 150$

Response time for process 3: $100 - 50 = 50$

Round Robin (Time slice = 100)




Wait time for process 1: $0 + (200 - 100) + (500 - 300) + (700 - 600) = 400$

Wait time for process 2: $(300 - 150) + (600 - 400) + (800 - 700) = 450$

Wait time for process 3: $(100 - 50) + (400 - 200) = 250$

Round Robin (Time slice = 100)



Turnaround time for process 1: $800 - 0 = 800$

Turnaround time for process 2: $900 - 150 = 750$



Turnaround time for process 3: $500 - 50 = 450$

FIFO vs. Round Robin

- With zero-cost context switch, is RR always better than FIFO?

FIFO vs. Round Robin

- Suppose we have three jobs of equal length



FIFO vs. Round Robin

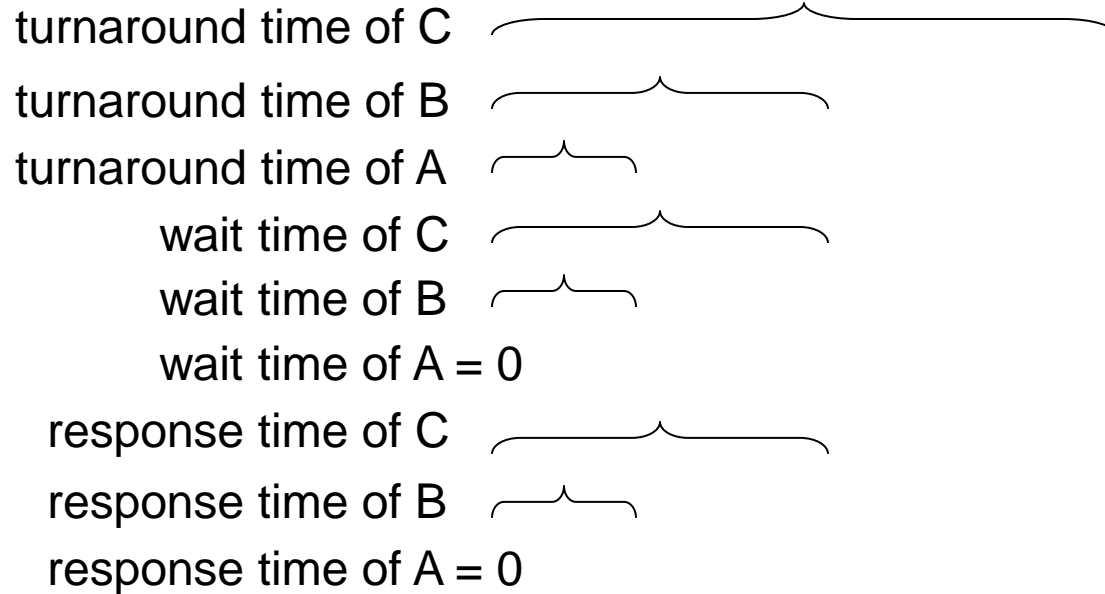
- Round Robin

- + Shorter response time
- + Fair sharing of CPU
- Not all jobs are preemptive
- Not good for jobs of the same length

Shortest Job First (SJF)

- ***SJF*** runs whatever job puts the least demand on the CPU, also known as ***STCF (shortest time to completion first)***
 - + Provably optimal
 - + Great for short jobs
 - + Small degradation for long jobs
- Real life example: supermarket express checkouts

SJF Illustrated



Shortest Remaining Time First (SRTF)

- **SRTF**: a preemptive version of SJF
 - If a job arrives with a shorter time to completion, SRTF preempts the CPU for the new job
 - Also known as **SRTCF** (*shortest remaining time to completion first*)
 - Generally used as the base case for comparisons


SJF and SRTF vs. FIFO and Round Robin

- If all jobs are the same length, SJF \rightarrow FIFO
 - FIFO is the best you can do
- If jobs have varying length
 - Short jobs do not get stuck behind long jobs under SRTF

A More Complicated Scenario (Arrival Times = 0)

- Process A (6 units of CPU request)

- 100% CPU
- 0% I/O



Time

- Process B (6 units of CPU request)


- 100% CPU
- 0% I/O



Time

- Process C (infinite loop)

- 33% CPU
- 67% I/O

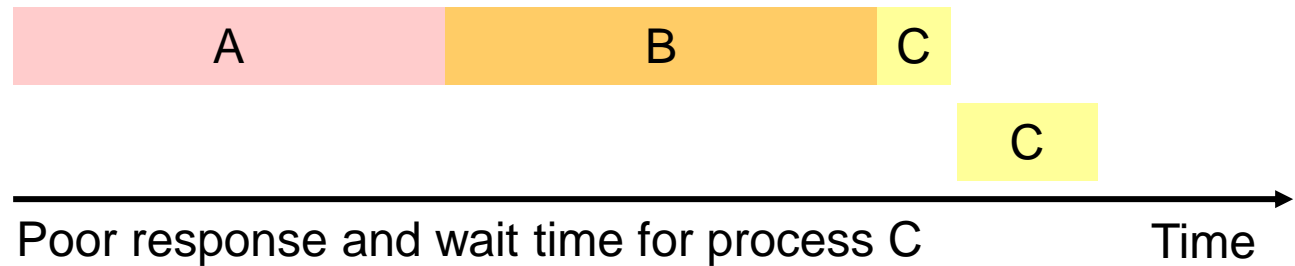


Time

A More Complicated Scenario


- FIFO

- CPU
- I/O



- Round Robin with time slice = 3 units


- CPU
- I/O



A More Complicated Scenario


- Round Robin with time slice = 1 unit

- CPU
- I/O



- SRTCF


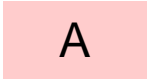

- CPU
- I/O




Drawbacks of Shortest Job First

- **Starvation**: constant arrivals of short jobs can keep long ones from running
- There is no way to know the completion time of jobs (most of the time)
 - Some solutions
 - Ask the user, who may not know any better
 - If a user cheats, the job is killed

Priority Scheduling (Multilevel Queues)

- **Priority scheduling**: The process with the highest priority runs first
- Priority 0:  C
- Priority 1:  A
- Priority 2:  B
- Assume that low numbers represent high priority



Priority Scheduling

+ Generalization of SJF

- With SJF, higher priority is inversely proportionally to requested_CPU_time

- Starvation

Multilevel Feedback Queues

- *Multilevel feedback queues* use multiple queues with different priorities
 - Round robin at each priority level
 - Run highest priority jobs first
 - Once those finish, run next highest priority, etc
 - Jobs start in the highest priority queue
 - If time slice expires, drop the job by one level
 - If time slice does not expire, push the job up by one level

Multilevel Feedback Queues

- Priority 0 (time slice = 1):
- Priority 1 (time slice = 2):
- Priority 2 (time slice = 4):

time = 0




Time



Multilevel Feedback Queues


- Priority 0 (time slice = 1):
- Priority 1 (time slice = 2):
- Priority 2 (time slice = 4):



Multilevel Feedback Queues

- Priority 0 (time slice = 1):
- Priority 1 (time slice = 2):
- Priority 2 (time slice = 4):

time = 2



Multilevel Feedback Queues

- Priority 0 (time slice = 1):
- Priority 1 (time slice = 2):
- Priority 2 (time slice = 4):

time = 3



Multilevel Feedback Queues

- Priority 0 (time slice = 1):
- Priority 1 (time slice = 2):
- Priority 2 (time slice = 4):

time = 3




suppose process A is blocked on an I/O



Multilevel Feedback Queues

- Priority 0 (time slice = 1):
- Priority 1 (time slice = 2):
- Priority 2 (time slice = 4):




suppose process A is blocked on an I/O



Multilevel Feedback Queues

- Priority 0 (time slice = 1):
- Priority 1 (time slice = 2):
- Priority 2 (time slice = 4):

time = 5



suppose process A is returned from an I/O



Multilevel Feedback Queues

- Priority 0 (time slice = 1):
- Priority 1 (time slice = 2):
- Priority 2 (time slice = 4):

time = 6

C




Multilevel Feedback Queues

- Priority 0 (time slice = 1):
- Priority 1 (time slice = 2):
- Priority 2 (time slice = 4):

time = 8


C



Multilevel Feedback Queues

- Priority 0 (time slice = 1):
- Priority 1 (time slice = 2):
- Priority 2 (time slice = 4):

time = 9



Multilevel Feedback Queues

- Approximates SRTF
 - A CPU-bound job drops like a rock
 - I/O-bound jobs stay near the top
 - Still unfair for long running jobs
 - Counter-measure: *Aging*
 - Increase the priority of long running jobs if they are not serviced for a period of time
 - Tricky to tune aging

Lottery Scheduling

- *Lottery scheduling* is an adaptive scheduling approach to address the fairness problem
 - Each process owns some tickets
 - On each time slice, a ticket is randomly picked
 - On average, the allocated CPU time is proportional to the number of tickets given to each job

Lottery Scheduling

- To approximate SJF, short jobs get more tickets
- To avoid starvation, each job gets at least one ticket

Lottery Scheduling Example

- short jobs: 10 tickets each
- long jobs: 1 ticket each

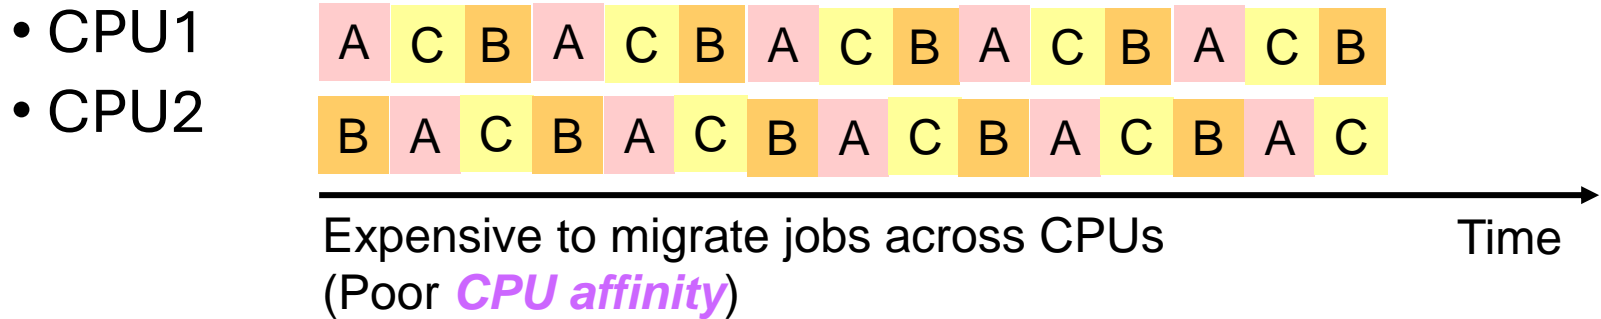
# short jobs/# long jobs	% of CPU for each short job	% of CPU for each long job

Pros and Cons of Lottery Scheduling


- + Good for coordinating computers with different computing power
- + Good for controlling the schedules for child processes
- Not as good for real-time systems

Multicore Scheduling

- Single-queue multiprocessor scheduling



- Another SQMS

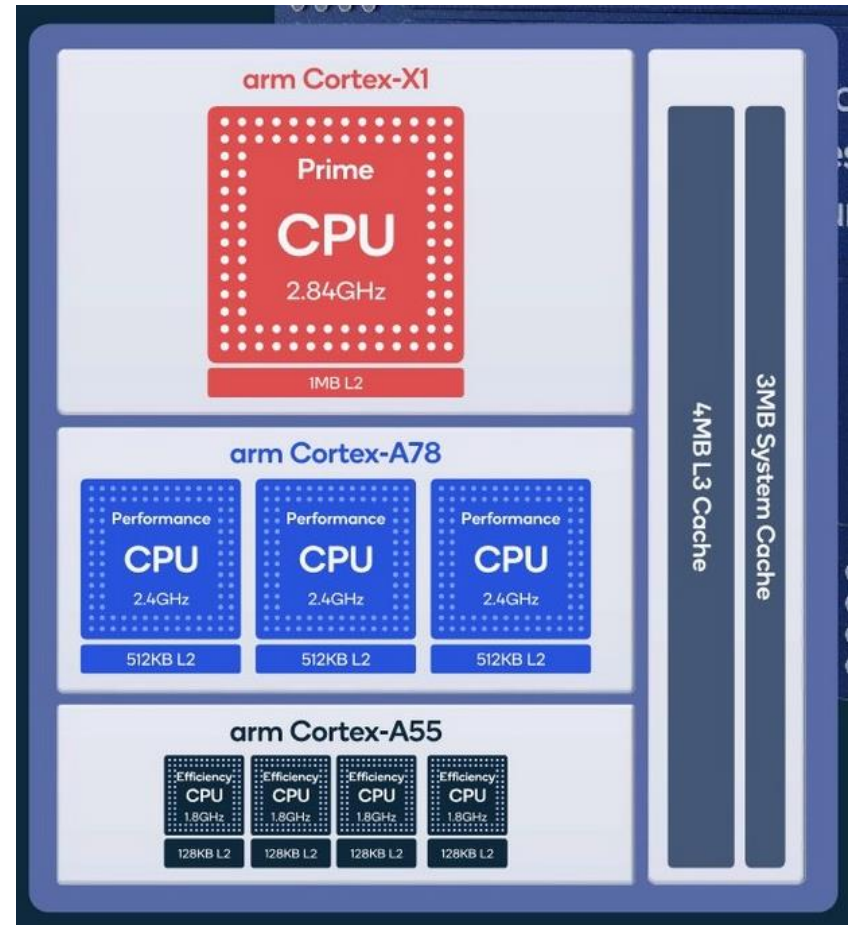


More Schedulers

- Multi-queue scheduling
- $O(1)$ scheduler
- Completely Fair Scheduler (CFS)

Real World

- Big.LITTLE
 - Snapdragon 888
- Others
 - Distance
 - Power
 - ...



Takeaways

- OS boot sequence
- Process, thread, and address space
- CPU scheduling policies