

# Cooperating Threads and Synchronization

Xin Liu

xl24j@fsu.edu

COP 4610 Operating Systems

# Teaching Assistants

- **Teaching Assistant 1: Michael Nguyen**
- **Office Hours: Friday, 1:00 PM – 4:00 PM, Major Lab.**
- **Teaching Assistant 2: Bing Jiao**
- **Office Hours: Tuesday, 2:00 PM – 5:00 PM, Major Lab.**

# Outline

- Independent Threads
- Cooperating Threads
- Race Condition
- Loss of Atomicity
- Synchronization

# Process

- An address space + at least one thread of execution
  - Address space offers protection among processes
  - Threads offer concurrency
- A fundamental unit of computation

# This Is How Simple Operating Systems Are!

CPU Reset → Firmware (BIOS/UEFI) → Boot Loader (MBR, LILO/GRUB)

↓  
Kernel\_start()

Use **ls /sbin/init -l** to check if systemd is being used

Read Latest Linux Kernel Code:

<https://elixir.bootlin.com/linux/v6.10.9/source/init/main.c#L1523>

↓  
Process 1

↓  
Application Program (state machine)  
+

system call

↓  
Process  
Management

↓  
Memory  
Management


↓  
File  
Management

- Process Management
  - fork, exec, and exit
- Memory Management
  - mmap –virtual address space
- File management
  - open, close, read, write
  - mkdir, link, unlink

You can use system call to create the world!

# Thread

- A sequential execution stream
  - The smallest CPU scheduling unit
  - Can be programmed as if it owns the entire CPU
    - Implication: an infinite loop within a thread won't halt the system
  - Illusion of multiple CPUs on a single-CPU machine




# Concurrency

- Allows multiple applications to run at the same time
  - Analogy: juggling



# User and Kernel Threads





# Independent Threads

- No states shared with other threads
- Deterministic computation
  - Output depends solely on the input
  - Same input always produces the same output
- Reproducible
  - Output does not depend on the order and timing of other threads
  - Scheduling order does not matter

# Pthreads

- May be provided either as user-level or kernel-level
- A POSIX standard (IEEE 1003.1c) API for thread creation and synchronization
- ***Specification***, not ***implementation***
- API specifies behavior of the thread library, implementation is up to development of the library
- Common in UNIX operating systems (Linux & Mac OS X)
- `pthread.h` defines the interface for pthreads

# Pthreads Example (Cont.)

```
#include <pthread.h>
```

```
#include <stdio.h>
```

```
#include <stdlib.h>
```

```
void *Ta(void *arg) {
```

```
    while (1) {
```

```
        printf("a");
```

```
    }
```

```
    return NULL;
```

```
}
```

```
void *Tb(void *arg) {
```

```
    while (1) {
```

```
        printf("b");
```

```
    }
```

```
    return NULL;
```

```
}
```

```
int main() {
```

```
    pthread_t thread1, thread2;
```

```
    if (pthread_create(&thread1, NULL, Ta, NULL) != 0) {
```

```
        fprintf(stderr, "Failed to create thread1\n");
```

```
        exit(1);
```

```
    }
```

```
    if (pthread_create(&thread2, NULL, Tb, NULL) != 0) {
```

```
        fprintf(stderr, "Failed to create thread2\n");
```

```
        exit(1);
```

```
    }
```

```
    pthread_join(thread1, NULL);
```

```
    pthread_join(thread2, NULL);
```

```
    return 0;
```

```
}
```

# Pthreads Example (Cont.)

- The program utilizing pthread.h can be written to take advantage of multiple processors!
- The operating system will automatically place threads on different processors.
- When running in the background, you can observe the CPU usage exceeding 100%.

# Concurrent Programming in HPC

- The World's Most Expensive Sofa
  - The First Supercomputer (1976)
  - Single-processor system
  - 138 million FLOPs (Floating Point Operations per Second)
    - 40 times faster than IBM 370 at the time
    - Slightly better than embedded chips today
  - Processed large data sets with one instruction



# Features of HPC


“A technology that harnesses the power of supercomputers or computer clusters to solve complex problems requiring massive computation.”  
(IBM)

- Computation-Centric
  - System Simulation: Weather forecasting, energy, molecular biology
  - Artificial Intelligence: Neural network training
  - Mining: Pure hash computation
  - TOP 500 (<https://www.top500.org/>)
    - 1<sup>st</sup>: Frontier (8, 699,904 cores, 1206 PFLOS)

# Main Challenges of HPC

- How to Break Down Computation Tasks?
  - Computation Graphs need to be easy to parallelize
    - Task decomposition happens on two levels: machine and thread
  - Parallel and Distributed Computation: Numerical Methods
- How Do Threads Communicate?
  - Communication happens not only between nodes/threads but also with any shared memory access
  - MPI - "a specification for developers and users of message-passing libraries"
  - OpenMP - "multi-platform shared-memory parallel programming in C/C++ and Fortran"

# Example: Mandelbrot Set



$$z_{n+1} = z_n^2 + c$$

Each point in the Mandelbrot set iterates independently and is only influenced by its complex coordinate



# Cooperating Threads

- Shared states (address space -> memory)
  - The Root of All Evil
- Nondeterministic
  - Output depends on **input and other factors**
- Nonreproducible
  - Same input can produce **different outputs** in different runs
  - Influenced by factors such as thread scheduling, randomness, or external environment

# Example: Share Memory

```
#include <pthread.h>
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>

#define NTHREAD
int x = 0;
pthread_t threads[NTHREAD];

void *Thello(void *arg) {
    int id = *(int *)arg;
    usleep(id * 100000);
    printf("Hello from thread #%%c\\n",
"123456789ABCDEF"[x++]);
    return NULL;
}
```

```
int main() {
    int ids[NTHREAD];

    for (int i = 0; i < NTHREAD; i++) {
        ids[i] = i + 1;

        if (pthread_create(&threads[i], NULL, Thello,
&ids[i]) != 0) {
            fprintf(stderr, "Error creating thread %%d\\n", i);
            return 1;
        }
    }

    for (int i = 0; i < NTHREAD; i++) {
        pthread_join(threads[i], NULL);
    }

    return 0;
}
```

# Example: Share Memory

```
#include <pthread.h>
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>

#define NTHREAD
int x = 0;
pthread_t threads[NTHREAD];

void *Thello(void *arg) {
    int id = *(int *)arg;
    usleep(id * 100000);
    printf("Hello from thread #%%c\\n",
"123456789ABCDEF"[x++]);
    return NULL;
}
```

```
int main() {
    int ids[NTHREAD];

    for (int i = 0; i < NTHREAD; i++) {
        ids[i] = i + 1;

        if (pthread_create(&threads[i], NULL, Thello,
&ids[i]) != 0) {
            fprintf(stderr, "Error creating thread %%d\\n", i);
            return 1;
        }
    }

    for (int i = 0; i < NTHREAD; i++) {
        pthread_join(threads[i], NULL);
    }

    return 0;
}
```

# So, Why Allow Cooperating Threads?

# So, Why Allow Cooperating Threads?

- Shared resources
  - e.g., a single processor
- Speedup
  - Occurs when threads use different resources at the same times
- Modularity
  - An application can be decomposed into threads

# However, Something Terrifying is Approaching..

- In a multiprocessor system, threads may execute code simultaneously.
- What will happen if two threads execute `x++` at the same time?

# Atomic Operations

- Atomicity refers to an operation or a sequence of operations that either completes fully or does not execute at all, without being interrupted by other operations during execution.
- It guarantees the indivisibility of an operation, ensuring that it cannot be partially completed or interrupted by another thread or process
- Key Characteristics:
  - Indivisibility: Atomic operations cannot be divided; no other thread or process can see or modify the operation's intermediate state.
  - Completeness: An atomic operation either fully succeeds and completes all its tasks, or it does not execute at all. There is no partial state.
  - No Interference: In a multi-threaded environment, atomic operations are not affected or interrupted by other threads.

# Examples of Atomic Operations

- Simple Operation:

- On most processors, an operation like `int x = 1;` is atomic because it involves a single memory action that cannot be interrupted.

- Non-Atomic Operation:

- Operations like `x++` involve multiple steps:
  - Read the current value of `x`.
  - Increment `x` by 1.
  - Write the new value back to `x`.



# Some Concurrent Programs

- If threads share data, the final values are not as obvious

Thread A

$x = 1;$

$x = y + 1;$

Thread B

$y = 2;$

$y = y * 2;$

- What are the indivisible operations?

# All Possible Execution Orders

Thread A


$x = 1;$

$x = y + 1;$

Thread B

$y = 2;$

$y = y * 2;$



A decision tree

# All Possible Execution Orders

Thread A

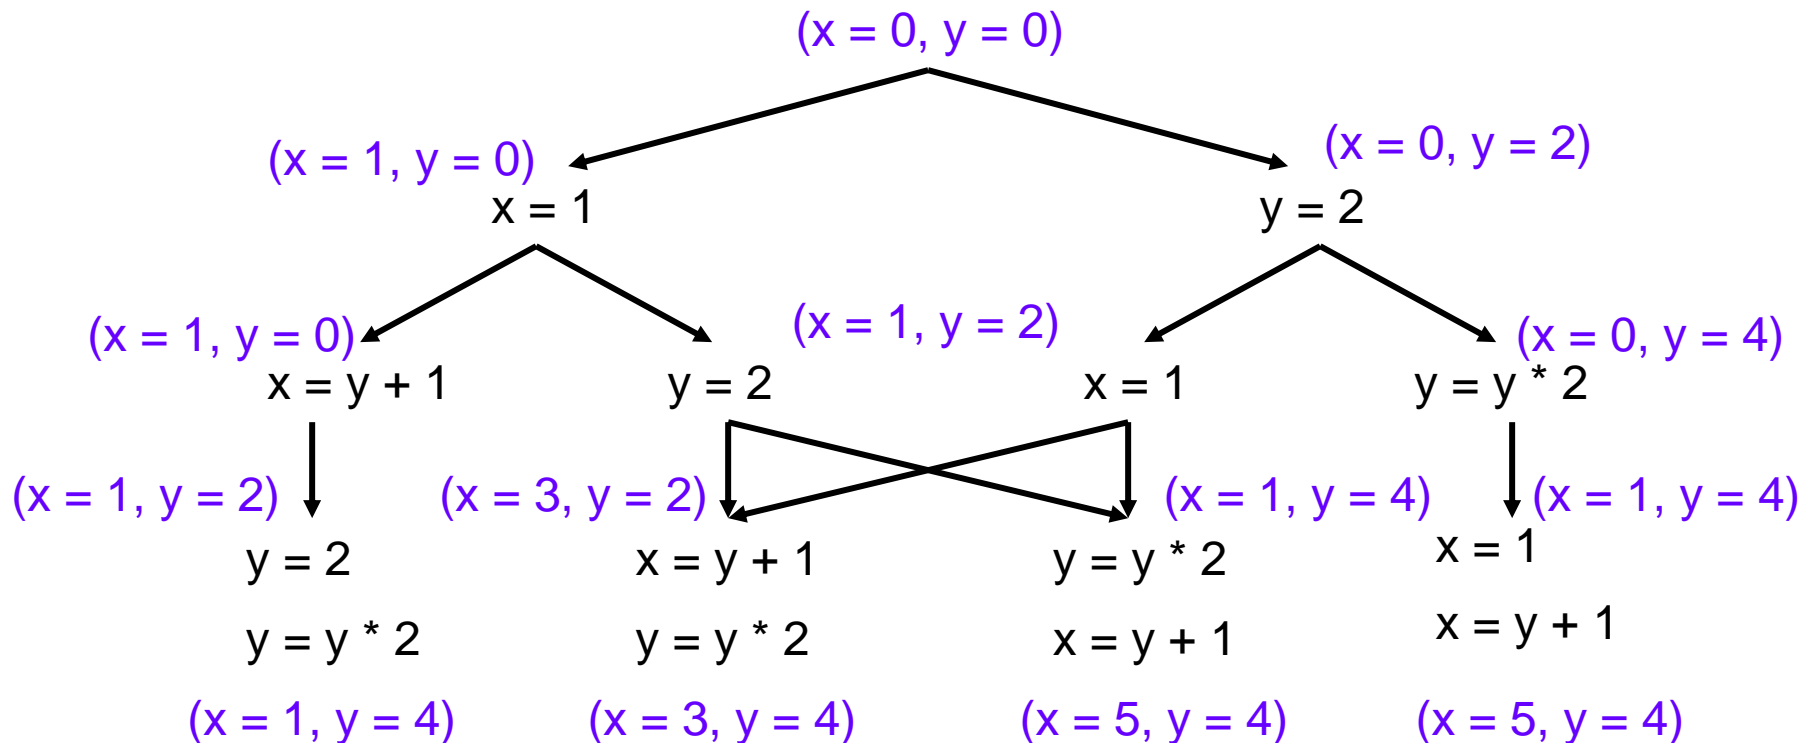
$x = 1;$

$x = y + 1;$

Thread B

$y = 2;$

$y = y * 2;$



# Race Condition

- Race conditions occur when threads share data, and their results depend on the timing of their execution.
- If we replace the `x++` operation in C with a single assembly instruction...
  - The atomicity of the operation is lost, as `x++` is not a single instruction but a series of steps (load, increment, store).
- Can this lead to race conditions in multithreaded environments?
- The atomicity of the operation is still lost.

# Loss of Atomicity in Modern Multiprocessor Systems

- The basic assumption that "a program (or even a single instruction) exclusively executes on the processor" no longer holds true in modern multiprocessor systems.
- Single Processor, Multithreading:
  - A thread may be interrupted and switched to another thread during execution.
- Multiprocessor, Multithreading:
  - Threads are executed truly in parallel.
- Historical Context (1960s):
  - There was a race to implement atomicity (mutual exclusion) in shared memory systems.
  - Almost all implementations were flawed until Dekker's Algorithm, which could only ensure mutual exclusion between two threads.

# Concurrent Programming in Data Centers

- Google Data Center



# Features of Data Center

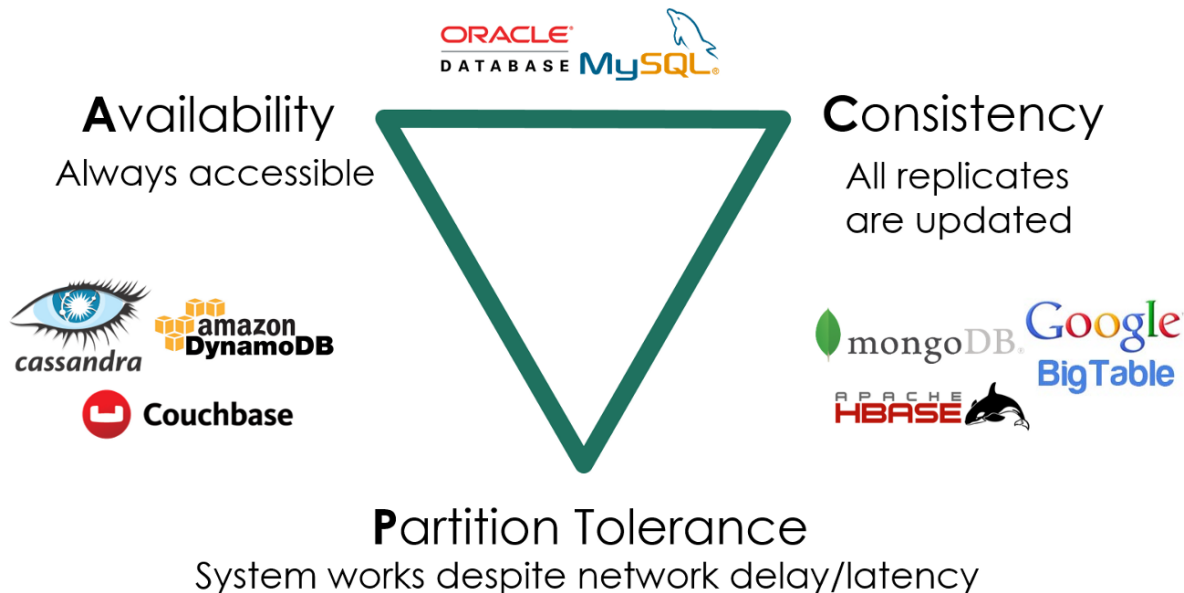
“A network of computing and storage resources that enable the delivery of shared applications and data.”

(CISCO)

- Data-Centric (Storage-Focused) Approach
  - Originated from internet search (Google), social networks (Facebook/Twitter)
  - Powers various internet applications: Gaming/Cloud Storage/ WeChat/Alipay/...
- The Importance of Algorithms/Systems for HPC and Data Centers
  - You manage 1,000,000 servers
  - A 1% improvement in an algorithm or implementation can save 10,000 servers

# Main Challenges of Data Center

- Highly Reliable and Low-Latency Data Access in Multi-Replica Systems
  - Serving massive, geographically distributed requests
  - Data must remain consistent (Consistency)
  - Services must always be available (Availability)
  - Must tolerate machine failures (Partition Tolerance)





# How to Maximize Parallel Request Handling with a Single Machine

- Key Metrics: QPS, Tail Latency, ...

- Tools We Have

- Threads

```
thread(start = true) {  
    println("${Thread.currentThread()} has run.")  
}
```

- Coroutines

- Multiple execution flows that can be paused/resumed (M2 - libco)
  - More lightweight than threads (no system calls, thus no OS state)

- GO

- Threads + Coroutines

# Concurrent Programming Around Us

- Web 2.0 Era (1999)
  - The Internet that connects people more closely.
  - “Users were encouraged to provide content, rather than just viewing it.”
  - You can even find some traces of “Web 3.0” / Metaverse.
- What Enabled Today’s Web 2.0?
  - Concurrent Programming in Browsers: Ajax (Asynchronous JavaScript + XML)
  - HTML (DOM Tree) + CSS
    - Represent everything you can see
  - JavaScript
    - Modify the page content
    - Connect local and server resources
  - You have the whole world at your fingertips!

# Features of Human-Computer Interaction

- As few concurrent tasks as possible, but just enough to meet requirements.
- One thread, a global event queue, and sequential execution.
- Run-to-completion: Each task runs to completion before the next one starts.

# Concurrent Programming - Real-world Applications

- High-Performance Computing
  - Focus: Task Decomposition
  - Pattern: Producer-Consumer
  - Technologies: MPI / OpenMP
- Data Centers
  - Focus: System Calls
  - Pattern: Threads-Coroutines
  - Technologies: Goroutine
- Human-Computer Interaction
  - Focus: Usability
  - Pattern: Event-Stream Graph
  - Technologies: Promise.

# Motivating Example: Too Much Milk

- Two robots are programmed to maintain the milk inventory at a store...
- They are not aware of each other's presence...



Robot: Dumb



Robot: Dumber

# Motivating Example: Too Much Milk

Dumb

10:00      Look into fridge:  
Out of milk

Dumber



# Motivating Example: Too Much Milk

Dumb

10:00 Look into fridge:

Out of milk

10:05 Head for the  
warehouse

Dumber



# Motivating Example: Too Much Milk

Dumb

10:05      Head for the  
warehouse

Dumber

10:10      Look into fridge:  
Out of milk





# Motivating Example: Too Much Milk

Dumb

Dumber

10:10      Look into fridge:  
             Out of milk

10:15      Head for the  
             warehouse



# Motivating Example: Too Much Milk

Dumb

10:20 Arrive with milk



Dumber

10:15 Head for the  
warehouse



# Motivating Example: Too Much Milk

Dumb

10:20 Arrive with milk



Dumber

10:15 Head for the  
warehouse



# Motivating Example: Too Much Milk

Dumb

10:20 Arrive with milk

10:25 Go party

Dumber



# Motivating Example: Too Much Milk

Dumb

10:20 Arrive with milk

10:25 Go party

Dumber

10:30 Arrive with milk:  
“Uh oh...”



# Definitions

- ***Synchronization***: uses atomic operations to ensure cooperation among threads
- ***Mutual exclusion***: ensures one thread can do something without the interference of other threads
- ***Critical section***: a piece of code that only one thread can execute at a time

# More on Critical Section

- A **lock** prevents a thread from doing something
  - A thread should lock before entering a critical section
  - A thread should unlock when leaving the critical section
  - A thread should wait if the critical section is locked
    - Synchronization often involves waiting

# Too Much Milk: Solution 1

- Two properties:
  - Only one robot will go get milk
  - Someone should go get the milk if needed
- Basic idea of solution 1
  - Leave a note (kind of like a lock)
  - Remove the note (kind of like a unlock)
  - Don't go get milk if the note is around (wait)



# Too Much Milk: Solution 1

```
if (no milk) {  
    if (no note) {  
        // leave a note;  
        // go get milk;  
        // remove the note;  
    }  
}
```

# Too Much Milk: Solution 1

Dumb

```
10:00 if (no milk) {
```



Dumber



# Too Much Milk: Solution 1

Dumb

```
10:00 if (no milk) {
```



Dumber

```
10:01 if (no milk) {
```



# Too Much Milk: Solution 1

Dumb

```
10:00 if (no milk) {
```



Dumber

```
10:01 if (no milk) {  
10:02     if (no note) {
```



# Too Much Milk: Solution 1

Dumb

```
10:00 if (no milk) {
```

```
10:03   if (no note) {
```



Dumber

```
10:01 if (no milk) {
```

```
10:02   if (no note) {
```



# Too Much Milk: Solution 1

Dumb

```
10:00 if (no milk) {
```

```
10:03   if (no note) {
```

```
10:04     // leave a note
```



Dumber

```
10:01 if (no milk) {
```

```
10:02   if (no note) {
```



# Too Much Milk: Solution 1

Dumb

```
10:03   if (no note) {  
10:04     // leave a note
```



Dumber

```
10:01  if (no milk) {  
10:02    if (no note) {  
  
10:05      // leave a note
```



# Too Much Milk: Solution 1

## Dumb

```
10:03  if (no note) {  
10:04      // leave a note  
  
10:06      // go get milk
```



## Dumber

```
10:02  if (no note) {  
  
10:05      // leave a note
```





# Too Much Milk: Solution 1

## Dumb

```
10:03  if (no note) {  
10:04      // leave a note  
  
10:06      // go get milk
```

## Dumber

```
10:05      // leave a note  
  
10:07      // go get milk
```



## Too Much Milk: Solution 2

- Okay...solution 1 does not work
- The notes are posted too late...
- What if both robots begin by leaving their own notes?

## Too Much Milk: Solution 2

```
// leave a note;  
if (no note from the other) {  
    if (no milk) {  
        // go get milk;  
    }  
}  
// remove the note;
```

# Too Much Milk: Solution 2

Dumb

10:00 // leave a note

Dumber



# Too Much Milk: Solution 2

Dumb

10:00 // leave a note



Dumber

10:01 // leave a note



# Too Much Milk: Solution 2

Dumb

```
10:00 // leave a note
```

```
10:02 if (no note from  
      Dumber) {...}
```



Dumber

```
10:01 // leave a note
```



# Too Much Milk: Solution 2

Dumb

```
10:00 // leave a note
```

```
10:02 if (no note from  
      Dumber) {...}
```



Dumber

```
10:01 // leave a note
```

```
10:03 if (no note from Dumb)  
      {...}
```



# Too Much Milk: Solution 2

Dumb

10:00 // leave a note

10:02 if (no note from  
Dumber) {...}

10:04 // remove the note



Dumber

10:01 // leave a note

10:03 if (no note from Dumb)  
{...}





# Too Much Milk: Solution 2

Dumb

```
10:00 // leave a note
```

```
10:02 if (no note from  
      Dumber) {...}
```

```
10:04 // remove the note
```



Dumber

```
10:01 // leave a note
```

```
10:03 if (no note from Dumb)  
      {...}
```



# Too Much Milk: Solution 2

Dumb

```
10:02 if (no note from  
      Dumber) {...}
```

```
10:04 // remove the note
```



Dumber

```
10:01 // leave a note
```

```
10:03 if (no note from Dumb)  
      {...}
```

```
10:05 // remove the note
```



# Too Much Milk: Solution 2

Dumb

```
10:02 if (no note from  
      Dumber) {...}
```

```
10:04 // remove the note
```



Dumber

```
10:01 // leave a note
```

```
10:03 if (no note from Dumb)  
      {...}
```

```
10:05 // remove the note
```



## Too Much Milk: Solution 2

- Solution 2 does not work
- The notes are found too late...
- What if both robots wait for the other to leave a note?

# Too Much Milk: Solution 3

Dumb

```
// leave Dumb's note
while (Dumber's note) { };
if (no milk) {
    // go get milk
}
// remove Dumb's note
```

Dumber

```
// leave Dumber's note
if (no Dumb's note) {
    if (no milk) {
        // go get milk
    }
}
// remove Dumber's note
```

# Too Much Milk Solution 3

- How do we verify the correctness of a solution?
- Test arbitrary interleaving of locking and checking locks
  - In this case, leaving notes and checking notes

# Dumber Challenges Dumb: Case 1

Dumb

```
// leave Dumb's note
while (Dumber's note) { };

if (no milk) {
    // go get milk
}

// remove Dumb's note
```

Dumber

```
// leave Dumber's note

if (no Dumb's note) {
}

// remove Dumber's note
```

Time



# Dumber Challenges Dumb: Case 2

Dumb

```
// leave Dumb's note

while (Dumber's note) { };

if (no milk) {
    // go get milk
}

// remove Dumb's note
```

Dumber

```
// leave Dumber's note

if (no Dumb's note) {
}

// remove Dumber's note
```

Time





# Dumber Challenges Dumb: Case 3

Dumb

```
// leave Dumb's note
```

```
while (Dumber's note) { };
```

```
if (no milk) {
```

```
    // go get milk
```

```
}
```

```
// remove Dumb's note
```

Dumber

```
// leave Dumber's note
```

```
if (no Dumb's note) {
```

```
}
```

```
// remove Dumber's note
```

Time  
↓

# Dumb Challenges Dumber: Case 1

Dumb

```
// leave Dumb's note  
while (Dumber's note) { };
```

```
if (no milk) {  
}  
// remove Dumb's note
```

Dumber

```
// leave Dumber's note  
if (no Dumb's note) {
```

```
    if (no milk) {  
        // go get milk  
    }
```

```
}  
// remove Dumber's note
```

Time  
↓

# Dumb Challenges Dumber: Case 2

Dumb

// leave Dumb's note

while (Dumber's note) { };

if (no milk) {

    // go get milk

}

// remove Dumb's note

Dumber

// leave Dumber's note

if (no Dumb's note) {

}

// remove Dumber's note

Time



# Dumb Challenges Dumber: Case 3

Dumb

```
// leave Dumb's note
while (Dumber's note) { };

if (no milk) {
    // go get milk
}

// remove Dumb's note
```

Dumber

```
// leave Dumber's note

if (no Dumb's note) {
}

// remove Dumber's note
```

Time



# Lessons Learned

- Although it works, Solution 3 is ugly
  - Difficult to verify correctness
  - Two threads have different code
    - Difficult to generalize to N threads
  - While Dumb is waiting, it consumes CPU time (*busy waiting*)
- More elegant with higher-level primitives  
lock→acquire();  
if (no milk) { // go get milk }  
lock→release();

# Takeaways

- Independent Threads
- Cooperating Threads
- Race Condition
- Loss of Atomicity
- Synchronization